

Computer Vision

N -View Reconstruction

dsai.asia

Asia Data Science and Artificial Intelligence Master's Program



Co-funded by the
Erasmus+ Programme
of the European Union



Readings for these lecture notes:

- Hartley, R., and Zisserman, A. (2004), *Multiple View Geometry in Computer Vision*, Cambridge University Press, Chapters 18–19.
- Triggs, B., McLauchlan, P., Hartley, R., and Fitzgibbon, A. (1999), Bundle adjustment — A modern synthesis, *Vision Algorithms: Theory and Practice*, Springer-Verlag.
- Tomasi and Kanade, Sturm and Triggs, Pollefeys et al., Davison et al., Klein and Murray, Lourakis and Argyros, Kerl, Sturm, and Cremers.

These notes contain material © Hartley and Zisserman (2004) and others.

Outline

- 1 Introduction
- 2 Bundle adjustment
- 3 Factorization
- 4 Resectioning
- 5 Metric upgrade and auto-calibration

Introduction

Reconstruction from N views

We have seen two-view projective and metric reconstruction techniques. As we move to many views, however, what can we do?

In this part we consider how to obtain a sparse reconstruction given a **sequence** of images.

This is where Hartley and Zisserman's book becomes a little obsolete, as there have been many new developments since 2004.

Introduction

History

A brief tour of the history:

- 1 1970's: The term **bundle adjustment** emerges in the photogrammetry literature, referring to simultaneous optimization of parameters of a set of cameras and a set of points observed by those cameras.
- 2 1992: Tomasi and Kanade show how to use SVD to factor the observation matrix to estimate a sequence of cameras and collection of 3D points. Limited to affine cameras with no missing points.
- 3 1996: Sturm and Triggs show how to use iterative factorization to obtain a projective reconstruction. Other factorization methods refine the technique.
- 4 2004: Pollefeys et al. combine keyframe selection, SfM, BA, resectioning, loop closure, autocalibration, and 3D mesh techniques to obtain textured 3D models from videos obtained with hand-held cameras.

Most offline 3D reconstruction methods use a pipeline similar to Pollefeys' approach, with manual intervention.

More recently, SfM (computer vision) and SLAM (robotics) techniques are starting to converge.

- 1 2006: Davison et al. introduce the first real-time monocular SLAM method, called **MonoSLAM**.
- 2 2007: Klein and Murray introduce **PTAM** (Parallel Tracking and Mapping) aimed at augmented reality applications.
- 3 2009: Lourakis and Argyros introduce **SBA** (Sparse BA), an efficient open source bundle adjustment library, bringing fast BA to the masses.
- 4 2013: Kerl, Sturm, and Cremers introduce **DVO SLAM** (Dense Visual SLAM) for RGB-D cameras.

As compute power increases, we are seeing more incremental real time methods with excellent results.

- ❶ 2014: Engel, Schöps, and Cremers introduce **SVO**, a semi-direct method for monocular visual odometry.
- ❷ 2014: Engel, Stückler, and Cremers (2015) introduce LSD SLAM (Large-Scale Direct Monocular SLAM), the first dense monoSLAM method.
- ❸ 2015: Mur-Artal, Montiel, and Tardós introduce ORB-SLAM, the most robust feature-based monoSLAM method today, combining the basic approach of PTAM with ORB features, BA, and loop closure techniques.

LSD SLAM brings the idea of dense photometric alignment from RGBD to RGB, providing richer maps than ORB-SLAM at higher computational cost.

Today, work is continuing on improving the robustness and efficiency of visual SLAM systems.

Several methods now obtain real time results on smartphones or embedded platforms like the Odroid XU4 or NVidia TX2.

However, it is still very difficult for state-of-the-art methods to keep track of a set of features during rotations and fast relative motion.

Visual-inertial SLAM systems attempt to combine IMU readings (linear acceleration and rotational velocity) with vision.

Knowing approximately how the camera has moved since the last keyframe gives us a better idea of where to look for features in the next frame.

Some visual-inertial SLAM systems:

- Christian Forster's Ph.D. thesis (2016) demonstrates combination of SVO with IMU preintegration to achieve very accurate VISLAM. No open source implementation.
- 2016: Forster, Zhang, Gassner, Welberger, and Scaramuzza introduce **SVO-2**, a faster, more accurate version of SVO incorporating Forster's IMU priors. Implementation is commercial (no open source).
- Raúl Mur-Artal, and Juan D. Tardós (2017) introduce **VI-ORB**, a visual-inertial version of ORB-SLAM using Forster's IMU preintegration. The authors did not release an open source version, but there is a community developed version on github.
- Stefan Leutenegger, Simon Lynen, Michael Bosse, Roland Siegwart and Paul Timothy Furgale (2015) introduce **OKVIS**, "Keyframe-based visual-inertial odometry using nonlinear optimization." Open-source version maintained by the author available on github.

Introduction

History

Interesting project for this class: VI-ORB or OKVIS experiments on Pixhawk or Android...

Outline

- 1 Introduction
- 2 Bundle adjustment**
- 3 Factorization
- 4 Resectioning
- 5 Metric upgrade and auto-calibration

Bundle adjustment

The idea

Given a set of unknown 3D points X_j viewed by a set of cameras with unknown projection matrices P^i at image points x_j^i , we seek to find the camera matrices P^i and 3D points X_j minimizing the reprojection error

$$\min_{\hat{P}^i, \hat{X}_j} \sum_{ij} d(\hat{P}^i \hat{X}_j, x_j^i)^2$$

Iterative minimization of this cost function is called **bundle adjustment** because it involves adjusting the bundle of rays between each camera center and the set of 3D points.

As formulated above, we need outlier removal prior to nonlinear least squares optimization. However, Triggs, McLauchlan, Hartley, and Fitzgibbon (1999) argue for a formulation with a more general cost function allowing robust estimation with the outliers included.

Bundle adjustment

Problems with bundle adjustment

There are two big **problems** with bundle adjustment:

- It needs a good **starting point** to begin optimization.
- There can be a **large number of parameters** involved in the minimization. For n points viewed by m cameras we have $3n + 11m$ parameters. This makes the matrices used by Levenberg-Marquardt prohibitively large.¹

¹Remember the linear system we have to solve on each iteration of LM?

$$\left[J_f^T(P) J_f(P) + \mu I \right] \delta P = -J_f^T(P) f(P).$$

Bundle adjustment

Solutions to the problems with bundle adjustment

Solutions of the first problem (the initial solution) generally involve linear methods such as **factorization**.

Solutions to the second problem (large parameter matrix):

- Reduce n and/or m , by using a subset of the points or partitioning the views [**suboptimal**].
- Interleave estimation of camera matrices with estimation of 3D points [guaranteed to converge but **slow**].
- Use a sparse minimization routine:
 - Download the sba (sparse bundle adjustment) open source software from <http://www.ics.forth.gr/~lourakis/sba/>
 - (For Debian) install the liblapack-dev package, and add `-llapack` to your gcc command line.
 - Call `sba_motstr_levmar()` to estimate motion (P_i 's) and structure (X).

Outline

- 1 Introduction
- 2 Bundle adjustment
- 3 Factorization**
- 4 Resectioning
- 5 Metric upgrade and auto-calibration

Factorization

Affine factorization

Tomasi and Kanade's algorithm is a maximum likelihood reconstruction in the case of **affine** cameras.

It requires that **all points** be seen in **all views**.

For affine cameras we write $\mathbf{x} = (x, y)^T$ and $\mathbf{X} = (X, Y, Z)^T$. Then we have the projection equation

$$\mathbf{x} = \mathbf{M}\mathbf{X} + \mathbf{t}$$

We seek cameras $\{\mathbf{M}^i, \mathbf{t}^i\}$ and 3D points \mathbf{X}_j such that the distance between estimated and predicted image points is minimized:

$$\min_{\mathbf{M}^i, \mathbf{t}^i, \mathbf{X}_j} \sum_{ij} (\mathbf{x}_j^i - \hat{\mathbf{x}}_j^i)^2 = \min_{\mathbf{M}^i, \mathbf{t}^i, \mathbf{X}_j} \sum_{ij} (\mathbf{x}_j^i - (\mathbf{M}^i \mathbf{X}_j + \mathbf{t}^i))^2.$$

Factorization

Affine factorization

If we choose the centroid of the points to be the origin of the coordinate system, then we can estimate the **translation vectors** t^i easily.

Under affine projection, the origin of the coordinate system is projected to $(0, 0)^T$ in the image. This means that t needs to translate the projection of the origin to the mean of the observed image points:

$$t^i = \frac{1}{n} \sum_j x_j^i.$$

To simplify the remaining calculations, we set $t^i = 0$ and adjust the points x_j^i accordingly.

Factorization

Affine factorization

Now we arrange the adjusted image points in the $2m \times n$ **measurement matrix**

$$W = \begin{bmatrix} x_1^1 & x_2^1 & \cdots & x_n^1 \\ x_1^2 & x_2^2 & \cdots & x_n^2 \\ \cdots & \cdots & \ddots & \cdots \\ x_1^m & x_2^m & \cdots & x_n^m \end{bmatrix}.$$

We want to find M^i and X_j such that

$$W = \begin{bmatrix} M^1 \\ M^2 \\ \cdots \\ M^m \end{bmatrix} [X_1 \quad X_2 \quad \cdots \quad X_n].$$

Factorization

Affine factorization

Since M and X are rank 3, their product is rank 3.

Since in general W will not be rank 3 due to measurement noise, we replace it with the best possible reconstruction, the matrix \hat{W} which is rank 3 and closest to W in Frobenius norm.

Such a matrix \hat{W} can be computed easily using the SVD $UDV^T = W$ by truncating U and V to three columns to get \hat{U} and \hat{V} , then truncating D to a 3×3 matrix \hat{D} , then finally letting $\hat{W} = \hat{U}\hat{D}\hat{V}^T$.

Since \hat{W} minimizes $\|W - \hat{W}\|_F$, $\hat{M} = \hat{U}$ and $\hat{X} = \hat{V}^T$ is a maximum likelihood reconstruction.

So we see that a straightforward application of the SVD gives us an **optimal reconstruction** in the case of **affine cameras**.

Factorization

Projective factorization

The affine factorization method doesn't work for **projective reconstruction**, so any serious projective distortion will introduce error into the reconstruction.

Sturm and Triggs (1996), however, pointed out that if we knew the **projective depth** of each point, then the structure and motion (camera matrices) could be estimated correctly by factorization.

We have $x_j^i = P^i X_j$. With a projective camera we have homogeneous points and there is an implicit constant factor which we can make explicit as $\lambda_j^i x_j^i = P^i X_j$ with $x_j^i = (x_j^i, y_j^i, 1)^T$.

Factorization

Projective factorization

If we write the depths explicitly and all points are visible in all images, we can write the problem in terms of a **scaled measurement matrix**:

$$W = \begin{bmatrix} \lambda_1^1 x_1^1 & \lambda_2^1 x_2^1 & \cdots & \lambda_n^1 x_n^1 \\ \lambda_1^2 x_1^2 & \lambda_2^2 x_2^2 & \cdots & \lambda_n^2 x_n^2 \\ \vdots & \vdots & \ddots & \vdots \\ \lambda_1^m x_1^m & \lambda_2^m x_2^m & \cdots & \lambda_n^m x_n^m \end{bmatrix} = \begin{bmatrix} p^1 \\ p^2 \\ \vdots \\ p^m \end{bmatrix} \begin{bmatrix} X_1 & X_2 & \cdots & X_n \end{bmatrix}$$

Since P and X are each rank 4, W will be rank 4 and the **rank 4 factorization** from the SVD will give a valid P and X .

Factorization

Projective factorization

The key to projective factorization is how to choose the **projective depths**² λ_j^i ?

- We could use another reconstruction method to estimate λ_j^i .
- We can also initialize them arbitrarily, e.g. $\lambda_j^i = 1$, then **interleave** factorization with depth estimation.

Though there is no guarantee that the iterative method will converge to a global minimum, it is widely used in practice.

²The λ_j^i are called projective depths because in a Euclidean frame they would be the lengths of the projections of the scene points onto the cameras' principal axes.

Factorization

Projective factorization

The algorithm works well in practice, but it is important to know **what cost function** is it minimizing.

By using rank 4 decomposition, it turns out the algorithm is minimizing

$$\|W - \hat{W}\|_F^2 = \sum_{ij} \|\lambda_j^i x_j^i - \hat{\lambda}_j^i \hat{x}_j^i\|^2 = \sum_{ij} (\lambda_j^i x_j^i - \hat{\lambda}_j^i \hat{x}_j^i)^2 + (\lambda_j^i y_j^i - \hat{\lambda}_j^i \hat{y}_j^i)^2 + (\lambda_j^i - \hat{\lambda}_j^i)^2$$

When the λ_j^i are close to each other, we see that the cost function is an **approximation to a scaled geometric distance**.

Factorization

Projective factorization

Because the projective reconstruction cost function involves λ_j^i , and the projective factorization method minimizes geometric error when $\forall i, j, \lambda_j^i = 1$, we would like to keep the λ_j^i as close to 1 as possible.

If we scale P and X , we obtain

$$(\alpha^i \beta_j \lambda_j^i) x_j^i = (\alpha^i P^i)(\beta_j X_j)$$

which means we can replace the projective depths by multiplying the i th row or j th column of W by an arbitrary factor.

One normalization method producing good results in practice is to renormalize, on every iteration, the rows and columns of W so they have unit norm.

As always, the image points should be normalized by isotropic scaling before beginning.

Factorization

Projective factorization

Projective factorization: Objective

Given a set of n image points seen in m views:

$$\mathbf{x}_j^i; i = 1, \dots, m, \quad j = 1, \dots, n$$

compute a projective reconstruction.

Factorization

Projective factorization

Projective factorization: Algorithm

- (i) Normalize the image data using isotropic scaling.
- (ii) Set projective depths $\lambda_j^i = 1$ or use some other method to estimate them.
- (iii) Normalize the depths λ_j^i by multiplying rows and columns by constant factors. One way is to do a pass setting the norm of each row to 1 then a pass setting the norm of each column to 1.
- (iv) Form the $3m \times n$ scaled measurement matrix W and obtain P^i and X_j from the rank 4 factorization of W .
- (v) Reproject the points into each image to obtain new estimates of the depths and repeat from step (iii) until convergence.

Factorization

Factorization approaches to N -view reconstruction

Factorization methods for reconstruction have received a great deal of attention over the last 10 years.

See Tang and Hung (2006) for a method based on the projective factorization idea that **allows missing points** and is also **guaranteed to converge** to a local minimum.

Outline

- 1 Introduction
- 2 Bundle adjustment
- 3 Factorization
- 4 Resectioning**
- 5 Metric upgrade and auto-calibration

Resectioning

Motivation

The factorization methods just outlined are **batch** methods.

They cannot be used to perform **on-line** 3D estimation without modification.

The **resectioning method** first gets an initial 3D reconstruction using, for example, the SVD of E , then repeats the following steps for frame i :

- Get correspondences between already-estimated 3D points $\{X_j\}_{j \in 1 \dots n}$ and new 2D points $\{x_k\}_{k \in 1 \dots m}$.
- Use the correspondences to estimate P^i . This is called **resectioning**.
- Use triangulation to estimate new 3D points using correspondences between frame i and frame $i - 1$.

Resectioning method

Challenges in resectioning

The main challenges in resectioning are **accumulated error**, **outliers**, and **degenerate conditions**.

To reduce accumulated error, we can periodically use bundle adjustment to find a globally consistent, minimum error configuration of the cameras and points.

To mitigate the effect of outliers, we use RANSAC or other robust estimators and outlier rejection.

To avoid degenerate conditions, we apply **keyframe selection**, in which we choose the images we use for reconstruction and resectioning specifically to avoid degenerate configurations of the cameras and points.

This leads to the general algorithm, on the next page.

Resectioning method

The algorithm

Reconstruct from an image sequence: Algorithm II

- (i) Compute **interest points** in each image using, e.g., SIFT.
- (ii) Extract **keyframes** from the image sequence in which significant motion separates successive keyframes.
- (iii) Obtain 2-frame reconstruction from keyframes 1 and 2.
- (iv) Perform **resectioning** on remaining images.
- (v) **Bundle adjust** the cameras and 3D structure for the complete keyframe sequence to obtain a maximum likelihood projective reconstruction.

Resectioning method

Keyframe selection

Keyframe selection **improves the runtime performance** of 3D reconstruction, helps improve **accuracy**, and helps avoid **degeneracy**.

Ahmed, Dailey, Landabaso, and Herrero (2010)³ experimented with a variety of criteria for robust key frame extraction.

Given the first keyframe, we apply **correspondence ratio** and **degeneracy avoidance** constraints to eliminate inappropriate candidate keyframes.

Then we select the best successor keyframe from the remaining candidate keyframe set by maximizing an objective function.

³Ahmed, M.T., Dailey, M.N., Landabaso, J.L., and Herrero, N. (2010), Robust key frame extraction for 3D reconstruction from video streams. In *International Conference on Computer Vision Theory and Applications (VISAPP)*.

Resectioning method

Keyframe selection

The **correspondence ratio** R_c is a proxy for the baseline:

$$R_c = \frac{T_c}{T_f},$$

where T_c is the number of inlier correspondences and T_f is total number of features found.

Low values of R_c indicate low overlap between two frames, in turn indicating long baselines.

Long baselines are desired for triangulation accuracy, but if the number of corresponding points is insufficient, camera motion estimation accuracy suffers.

Thus, the correspondence ratio is constrained to lie between an upper threshold T_1 and a lower threshold T_2 .

Resectioning method

Keyframe selection

To avoid degenerate cases, the **geometric robust information criterion (GRIC)** score is used to measure the goodness of fit for a homography or the fundamental matrix.

$$GRIC = \sum_i \rho(e_i^2) + \lambda_1 dn + \lambda_2 k,$$

where $i = 1 \dots n$, $\rho(e_i^2)$ is a robust function

$$\rho(e_i^2) = \min\left(\frac{e_i^2}{\sigma^2}, \lambda_3(r - d)\right)$$

of the residual e_i over the n correspondences, σ is the assumed standard deviation of the error, d is the model dimension, k is the model degrees of freedom, r is the dimension of the data, $\lambda_1 = \log(r)$, $\lambda_2 = \log(rn)$, and λ_3 limits the residual error.

The GRIC constraint is to **reject** any candidate keyframes for which the **homography** model has a **lower GRIC score** than the **fundamental matrix**.

Resectioning method

Keyframe selection

After filtering by correspondence ratio and degeneracy constraints, we select as the next keyframe the frame that maximizes certain selection criteria incorporating

- **Normalized GRIC difference:**

$$f_G(i, j) = \frac{GRIC_H(i, j) - GRIC_F(i, j)}{GRIC_H}.$$

The difference is maximized when the fundamental matrix model is much better than the homography model.

- **Point-to-epipolar line cost (PELC):**

$$f_{GP}(i, j) = w_G f_G(i, j) + w_P (\sigma - PELC(i, j)),$$

where i is the current keyframe and j is a candidate next keyframe. PELC tends to be high when correspondences are not accurate.

Resectioning method

Resectioning

The camera matrix P^i is computed from 3D-2D correspondences and the DLT algorithm.

In practice, 3D reconstruction using **long** video sequence suffers from **accumulated error**. Some issues are:

- **Matching existing 3D to 2D points** in an additional frame is not perfect. Sometimes, already-existing features are treated as new points. This arises when the camera moves back and forth or when a point becomes occluded and then visible again.
- Resectioning is highly sensitive to **outliers**. The camera estimate can be completely wrong if we estimate P^i in the presence of outliers.

Resectioning method

Example from AIT Vision and Graphics Lab

Atima Tharatipyakul worked on solving these problems by finding 2D-3D point pairs using multiple frames then using RANSAC in camera estimation to discard 2D-3D correspondences outliers.⁴


Next slide: example results from resectioning method in Atima's work, with 5 keyframes selected from a sequence of images containing pineapples.

3 of the keyframes are shown, along with camera estimates, point cloud estimates, and fruit ellipsoid estimates.

⁴Tharatipyakul, A., *3D Visualization from Video Sequence for Agricultural Field Inspection Robot*, Master's thesis, Asian Institute of Technology, 2011.

Resectioning method

Example from AIT Vision and Graphics Lab



Outline

- 1 Introduction
- 2 Bundle adjustment
- 3 Factorization
- 4 Resectioning
- 5 Metric upgrade and auto-calibration**

Metric upgrade and auto-calibration

Introduction

As we already know, if K is known we can directly obtain metric reconstructions. For two frames:

- For first pair of keyframes, estimate F .
- Calculate E from F and K .
- Factor E to obtain P' as described in H&Z Section 9.6.

For N frames, we factor E for the first pair of keyframes, then use incremental resectioning and bundle adjustment.

If K is **unknown** or **changing** during the image sequence, however, we need a method for **auto-calibration**.

Here we look at some of the techniques for auto-calibration from image correspondences over an image sequence.

Metric upgrade and auto-calibration

Idea of auto-calibration

The idea of auto-calibration:

- Obtain a projective factorization $W = PX$.
- Estimate a homography H such that $W = (PH)(H^{-1}X)$ is a metric reconstruction.
- $P'H$ is metric if it can be decomposed as $K^i[R^i \mid t^i]$ where the K^i are consistent with a priori constraints (the same across the image sequence, only focal length changing, etc.).

Metric upgrade and auto-calibration

Direct vs. stratified methods

There are two main approaches to auto-calibration for general motion:

- **Direct** estimation of H .
- **Stratified** reconstruction, beginning with an affine reconstruction (which identifies π_∞) followed by metric reconstruction.

Here we only consider direct methods, though stratified methods have some advantages, mainly that there is a linear solution for metric reconstruction once π_∞ is known (see text for details).

Metric upgrade and auto-calibration

Framework

We seek H such that $P_M^i = P^i H = K^i [R^i \mid t^i]$ for $i = 1, \dots, m$.

Since we don't care about the absolute frame, we assume $P^1 = [I \mid 0]$ and that therefore $P_M^1 = K^1 [I \mid 0]$.

In general, H takes the form

$$H = \begin{bmatrix} A & t \\ v^T & k \end{bmatrix}.$$

Since $P_M^1 = P^1 H = [I \mid 0]$, we can infer that $A = K^1$ and $t = 0$.

Since H is necessarily non-singular we can assume $k = 1$ to obtain

$$H = \begin{bmatrix} K^1 & 0 \\ v^T & 1 \end{bmatrix}.$$

Metric upgrade and auto-calibration

Framework

The plane at infinity π_∞ is $(0, 0, 0, 1)^T$ in an affine or metric frame, and H^{-T} is the point/plane transform from the metric frame to the projective frame. This means we can derive

$$\pi_\infty = \begin{pmatrix} -(K^1)^{-T} v \\ 1 \end{pmatrix}.$$

Finally we can write

$$H = \begin{bmatrix} K & 0 \\ -p^T K & 1 \end{bmatrix}, \pi_\infty = (p^T, 1)^T.$$

Metric upgrade and auto-calibration

Framework

For the rest of the cameras ($i = 2, \dots, m$), we write $P^i = [A^i \mid a^i]$.

Using $P_M^i = P^i H$, we can obtain

$$K^i R^i = (A^i - a^i p^T) K^1$$

and (since R^i is orthogonal),

$$K^i K^{iT} = (A^i - a^i p^T) K^1 K^{1T} (A^i - a^i p^T)^T.$$

These are the **basic auto-calibration equations**. If we know K^1 , p , and P^i , we can calculate P_M^i .

Metric upgrade and auto-calibration

The DIAC and the absolute dual quadric

The matrix $K^i K^{iT}$ is the **dual image of the absolute conic (DIAC)** ω^{*i} .

The DIAC is the projection of the **absolute dual quadric** Q_∞^* :

$$K^i K^{iT} = \omega^{*i} = P^i Q_\infty^* P^{iT}$$

If we know Q_∞^* , we can calculate K^i directly (by Cholesky decomposition of $P^i Q_\infty^* P^{iT}$).

Metric upgrade and auto-calibration

Framework

Important facts:

- The **absolute conic** Ω_∞ is a conic on π_∞ containing the intersection of all circles and spheres with π_∞ .
- The **absolute dual quadric** Q_∞^* is a rank 3 dual quadric whose envelope is the set of planes tangent to Ω_∞ .
- The absolute dual quadric is invariant under similarity transforms and is just $\text{diag}(1,1,1,0)$ in a metric frame.
- The absolute dual quadric's null space is π_∞ .

Metric upgrade and auto-calibration

Auto-calibration based on Q_{∞}^*

Auto-calibration based on Q_{∞}^* : Objective

Given a set of matched points across several views and constraints on the calibration matrices K^i , compute a metric reconstruction of the points and cameras.

Metric upgrade and auto-calibration

Auto-calibration based on Q_{∞}^*

Auto-calibration based on Q_{∞}^* : Algorithm

- (i) Compute a projective reconstruction from a set of views, resulting in cameras P^i and points X .
- (ii) Use $\omega^{*i} = P^i Q_{\infty}^* P^{iT}$ along with constraints to estimate Q_{∞}^* .
- (iii) Decompose Q_{∞}^* as $H \tilde{I} H^T$ where $\tilde{I} = \text{diag}(1, 1, 1, 0)$.
- (iv) Apply H^{-1} to X and H to P^i to obtain a metric reconstruction of the point and cameras.
- (v) Use iterative least squares to improve the solution.

Metric upgrade and auto-calibration

Auto-calibration based on Q_{∞}^*

As a nice example of linear constraints on Q_{∞}^* , see Pollefeys et al. (2004), Visual modeling with a hand-held camera, *IJCV* 59(3).

As a nice example of more sophisticated methods to enforce the rank 3, positive semidefiniteness, and “chirality” constraints on Q_{∞} , see Chandraker et al. (2007), Autocalibration via rank-constrained estimation of the absolute dual quadric, In *Proceedings of CVPR*.